ML Based Event Analytics and Participation Prediction System

Ayush Aman Gautam Sharma Shubham Gupta
ayush.aman11105@gmail.com gautamsharmaa4@gmail.com Shubu1178@gmail.com
B.Tech. (AI&DS), 4th Semester, Section S13

Under the Supervision of

Mr. Yug Lohchab
Department of Artificial Intelligence and Data Science
Dr. Akhilesh Das Gupta Institute of Professional Studies,New Delhi

***Abstract*—Predicting whether a registered participant will actually Being required to attend academic events is a common problem for most students. institutions leave unsolved, not because the data is unavailable but because no lightweight, interpretable to analyse event structures and it is the second iteration of the initial tool. This Paper presents EventIQ v2 ,a machine learning system designed specifically for the small, heterogeneous datasets that are typical of institutional settings. The system introduces a seven-term interaction characteristic layer – including `rating _x _org`, `buzz _x _rating` and `loyalty_score` – that encodes how event quality, organizer reputation, and social engagement combine to drive attendance, rather than treating each factor in isolation. A Gradient Boosting Classifier trained on 3,000 records with 18 engineered features achieves88% accuracy, AUC-ROC of 0.956, and F1 Score of 0.904. Five-fold cross-validation confirms that these results generalise stably (CV accuracy ≈88%), closing a 25-point Compared to our previous work, gap from the v1 prototype. An ablation study shows the contributions of each module. interaction layer alone contributes an AUC-ROC gain of 0.030. This product goes 0.030 above basic functions. The 0.030 product line consists of several models which can be installed. This is a Flask REST API that allows you to store and interrogate multiple streams of time-stamped and dated data. It is supplemented by a five-page Streamlit-based interactive web interface that includes per-prediction explanation charts, making model output legible to faculty coordinators and student organisers who have no machine learning background.**

***Index Terms*—Attendance Prediction, Gradient Boosting, Feature Engineering, Interaction Features, Academic Event Management, Interpretable Machine Learning, Data Scarcity, Ensemble Methods**

I. Introduction

Academic institutions invest considerable resources in organising events — technical workshops, guest lectures, cultural fests, and hackathons — that serve genuine educational and professional development purposes. One of the most persistent operational problems in running notably this changes the number of participants who engage with the festival. Few show up as compared to the number who registered. Over-provisioning venues and catering wastes institutional budgets; Under-provisioning creates a poor experience, so many people are not attending, and this is reflecting poorly on the organisation team..

I wish that data existed to explore this question better already exists. Every registration system generates records capturing who signed up, from where, when, at what income level, and with what prior attendance history. The problem is not data availability But unfortunately there is no pipeline to turn those records into widgets into useful pre-event forecasts. Commercial event analytics platforms are designed for large-scale deployments with hundreds of thousands of records and cloud infrastructure.Academic units frequently requires research data,information or analysis that a typical student organisation are unable to obtain.

Research Gap - Existing attendance prediction approaches have three key shortcomings in the academic institutional context: (i) they require data volumes that are not available in institutional settings; (ii) they treat input features independently, missing compound effects such as the combined influence of organiser reputation and event quality on attendance follow-through; and (iii) they produce outputs that faculty coordinators and student organisers cannot interpret or act on without technical training. EventIQ v2 addresses all three challenges mentioned above. gaps simultaneously.

Contributions - The specific contributions of this work are as follows.

- A seven term interaction feature framework, grounded in domain specific behavioural hypothesis that encodes compound relationships between quality ,reputation,social engagement and logistics.An ablation study confirms an independent AUC -ROC gain of ≈0.030 beyond base features.
- An empirical comparison of six classifiers on a common 18-feature set, with Gradient Boosting selected based on accuracy, AUC-ROC, and the practical trade-off between interpretability and deployment simplicity.
- A rigorous five-fold cross-validation evaluation on 3,000 records that closes the 20-point CV gap observed in the v1 prototype and confirms that the model generalises

reliably to unseen registrants.

- An end-to-end deployment as a Flask REST API and a five-page Streamlit dashboard with per-prediction driver charts that make model reasoning transparent to non-technical users.

II. Related Work and Research Gap

The problem of predicting whether someone will follow through on a commitment they have already made appears in many forms across different fields - demand forecasting in retail,patient no show prediction in healthcare and churn modeling in subscription services,all share the same underlying structure.In each case the key insight is that historical behavior is the most reliable available signal for future behavior. [3]. Academic event attendance prediction is a natural extension of this principle,though it has recieved far less dedicated attention in the research literature than these adjacent domains.

A. Taxonomy of Existing Approaches

Table I organises prior work into three methodological categories, with the key strengths and limitations of each approach noted.

TABLE I
Taxonomy of Participation Prediction Approaches

Category	Methods	Strengths	Limitations
Classical Statistical	LR, Ridge, Lasso [2]	Transparent, low data needs	Cannot model non-linear feature interactions
Ensemble / Tree	RF, XGBoost [1]	Strong accuracy, native feature importance	RF unstable on small n; XGB complex to tune
Social-Signal Augmented	RSVP + social features [7]	Captures online engagement signals	Requires large-scale social platform data

Classical statistical models. Linear and Logistic Regression are transparent and work reasonably well with limited data [2], which makes them appealing in institutional settings. Their fundamental limitation, however, is the assumption of feature independence. In practice, the effect of venue distance on attendance depends heavily on income . The challenge of capturing that feeling. relationship requires either explicit interaction terms or a model that can represent non-linearity natively [4]. Regularised variants such as Ridge and Lasso help with correlated features but do not address this core limitation.

Ensemble and tree-based methods. Breiman's Random Forest [1] trains Bootstrapping and using multiple decision trees on random subsets of the data and averages their outputs, reducing the variance that makes individual trees unreliable on small datasets. Its Gini-based feature importance scores are genuinely also really useful here: they give organisers a ranked list of what actually drives attendance rather than just a probability number. The practical limitation is that on small datasets (i.e. fewer than a few hundred samples, typically) the variance reduction from bagging is insufficient and the model can still overfit . XGBoost [8] builds trees sequentially to correct the errors of previous estimators and consistently produces stronger accuracy, but its sensitivity to hyperparameter tuning and its relative opacity make it harder to justify for a first deployment where building stakeholder trust matters as much as maximising accuracy. Gradient Boosting [9] provides a middle path: better calibration and accuracy than Random Forest, deployable at similar complexity, and with probability outputs that translate naturally into a planning dashboard.

Social-signal augmented models. Wang, Tang, and Liu [7] studied RSVP users behave differently on a large social event platform and therefore require customised search experiences and appropriate ranking strategies that social visibility — measured through shares and online mentions — had an independent positive effect on whether a confirmed RSVP converted to actual attendance, beyond what event type, timing, and location explained. That finding motivated the inclusion of `social_buzz` as a feature in EventIQ v2. The practical limitation of their approach is its dependence on social platform data at a scale that institutional events do not generate.

B. Identified Research Gap

No prior work in the attendance prediction literature operates simultaneously on small institutional datasets (under 5,000 records), models compound feature interactions explicitly rather than treating inputs independently, and delivers interpretable output through a lightweight deployment that does not require technical expertise to use. EventIQ v2 is designed to address all three of these requirements together.

III. Mathematical Formulation

A. Problem Statement

Let $\mathcal{D} = \{(\mathbf{x}_i, y_i)\}_{i=1}^{N}$ denote the training dataset of $N = 3{,}000$ records, where $\mathbf{x}_i \in \mathbb{R}^{18}$ is the engineered feature vector for registrant i and $y_i \in \{0, 1\}$ is the binary attendance label ($y = 1$ attended, $y = 0$ did not attend). The objective is to learn a classifier $f : \mathbb{R}^{18} \to [0, 1]$ that maps each feature vector to an attendance probability, with binary predictions obtained by applying a threshold of 0.5.

B. Interaction Feature Construction

The seven interaction terms are constructed as explicit multiplicative and ratio-based transformations of base features, each one encoding a specific hypothesis about An interaction exists between two or more independent variables that predict attendance. The highest-ranked feature, `loyalty_score`, is defined as:

$$\ell_i = \frac{p_i \cdot r_i \cdot s_i}{25} \tag{1}$$

where p_i is `previous_events`, r_i is `event_rating`, and s_i is `organizer_score`. The denominator scales the prod-

uct to a range comparable with other features. The income-distance interaction is defined as:

$$d_i^* = \frac{d_i}{(w_i/10000) + 1} \tag{2}$$

where d_i is `distance_km` and w_i is monthly income. This encodes the observation that Travel cost is not an absolute barrier but an income-barrier to traveling to wilderness areas ,the same distance means something very different at different income levels.

C. Gradient Boosting Objective

The Gradient Boosting Classifier minimises log-loss over $M = 300$ sequential estimators [9]. At each iteration m, a new weak learner $h_m(\mathbf{x})$ is fitted to the negative gradient of the current loss, and the ensemble is updated as:

$$F_m(\mathbf{x}) = F_{m-1}(\mathbf{x}) + \eta \cdot h_m(\mathbf{x}) \tag{3}$$

where $\eta = 0.05$ is the learning rate that controls how aggressively each new tree corrects its predecessor. The final model $F_M(\mathbf{x})$ is passed through a sigmoid to produce the attendance probability estimate.

D. Feature Importance via Gini Impurity

The importance of feature j is measured as the mean decrease in Gini impurity across all split nodes that use feature j, averaged over all M trees:

$$\mathrm{Imp}(j) = \frac{1}{M} \sum_{m=1}^{M} \sum_{\substack{t \in T_m \\ \mathrm{split}(t)=j}} \Delta G(t) \tag{4}$$

where $\Delta G(t)$ is the weighted Gini reduction at node t and T_m is the m-th tree. Importances are normalised to sum to 1.0 across all 18 features and reported in Table VI.

IV. Methodology

A. Dataset and Synthetic Augmentation

The v2 dataset contains 3,000 event registration records spanning six Indian cities — Pune, Delhi, Mumbai, Bangalore, Chennai, and Kolkata — and six event categories: Art, Business, Education, Music, Sports, Tech. The core of the dataset is built from real institutional records. Where certain category-city combinations had sparse coverage, additional records were generated by stratified bootstrap sampling from the empirical distributions compared to the real data observations. well-accepted approach for augmenting small institutional datasets [3];it preserves the statistical properties of the original data while ensuring that every combination has sufficient representation for reliable cross representation. Each record carries a binary attendance outcome.

The expansion from 120 records (v1) to 3,000 records was not arbitrary. With $N = 120$ and $K = 5$ folds, each cross-validation fold contains only 24 samples, which produces CV estimates with high variance and little reliability. With $N = 3{,}000$, each fold contains 600 samples, and empirical verification confirmed that CV accuracy across the five folds varies by under 2 percentage points — a stable and trustworthy generalisation estimate.

B. Preprocessing Pipeline

All preprocessing runs through a single shared module, `utils/preprocess.py`, which is used identically by both the training script and the inference endpoint. This design guarantees that the transformations applied to new inputs at prediction time are exactly the same as those applied during training, eliminating a common source of errors in deployed machine learning systems. The pipeline applies five steps in sequence: (1) deduplication; (2) median and mode imputation for missing values in numerical and categorical columns respectively; (3) IQR-based outlier capping at $[Q_1 - 1.5\sigma_{\mathrm{IQR}},\ Q_3 + 1.5\sigma_{\mathrm{IQR}}]$, which caps rather than drops extreme values to preserve every row; (4) `LabelEncoder` for `gender`, `location`, and `event_type`, with the fitted encoders serialised for consistent inference-time mapping; and (5) `StandardScaler` normalisation of all continuous features to zero mean and unit variance.

C. Feature Engineering

The key insight behind v2 is that individual features rarely determine attendance in isolation. A high event rating means more when the organiser also has a strong reputation. Social buzz matters more when the underlying event quality justifies it. Distance is a bigger barrier for lower-income participants than for higher-income ones. The 11-feature set from v1 treated everything independently and could not represent any of these relationships. V2 builds 18 features in total, including 7 interaction terms that encode these compound effects directly. Table II lists all 18.

`organizer_score` (1.0–5.0) captures the historical reputation of the event organiser, derived from post-event participant feedback on previous events they have run. The hypothesis is straightforward: people are more likely to follow through on a registration when they trust the organiser to deliver what was promised. `social_buzz` (0–100) aggregates social media engagement signals — mentions, shares, and reactions — around the event. In the dataset, events with a buzz score above 70 show a statistically distinct attendance conversion pattern compared to quieter events, consistent with findings in the RSVP prediction literature [7].

D. Model Selection and Configuration

Six classifiers were evaluated on the same 18-feature set under an identical 80/20 stratified split with five-fold cross-validation. Gradient Boosting was selected on the basis of four criteria: the highest AUC-ROC (0.956) of any model tested, well-calibrated probability outputs that translate directly into the dashboard probability gauge, a compact serialised size of approximately 1.2 MB that loads in under a second, and native Gini-based feature importance scores. Table III documents the final hyperparameter configuration.

E. Evaluation Protocol

Four metrics are reported across a consistent protocol. **Accuracy** gives overall correctness on the held-out test set. **AUC-ROC** measures the classifier's ability to rank a true attendee

TABLE II
FULL FEATURE SET — EVENTIQ V2 (18 FEATURES)

Feature	Type	Description
Base demographic / event (8)		
`age`	Continuous	Registrant age
`gender`	Categorical	Encoded (3 classes)
`location`	Categorical	City (6 classes)
`event_type`	Categorical	Category (6 classes)
`income`	Continuous	Monthly income (INR)
`event_rating`	Continuous	Past event rating (1–5)
`distance_km`	Continuous	Venue distance (km)
`previous_events`	Continuous	Prior events attended
Novel domain features (2)		
`organizer_score`	Continuous	Reputation score (1–5)
`social_buzz`	Continuous	Online engagement (0–100)
Temporal features (3)		
`days_to_event`	Derived	Registration lead time (days)
`event_month`	Derived	Calendar month of event
`reg_month`	Derived	Calendar month of sign-up
Interaction features (7)		
`rating_x_org`	Interaction	$r_i \times s_i$
`buzz_x_rating`	Interaction	$b_i \times r_i/100$
`prev_x_org`	Interaction	$p_i \times s_i$
`income_per_km`	Interaction	$w_i/(d_i+1)$
`rating_x_prev`	Interaction	$r_i \times p_i$
`dist_income`	Interaction	$d_i/(w_i/10000+1)$
`loyalty_score`	Composite	$p_i r_i s_i/25$ (Eq. 1)

TABLE III
GRADIENT BOOSTING HYPERPARAMETER CONFIGURATION

Parameter	Value	Rationale
`n_estimators`	300	Stable ensemble convergence
`learning_rate`	0.05	Conservative shrinkage; reduces overfitting
`max_depth`	5	Controls individual tree complexity
`subsample`	0.8	Stochastic boosting reduces variance
`min_samples_split`	10	Prevents leaf-level noise fitting
`random_state`	42	Full reproducibility

above a true non-attendee across all probability thresholds, which is particularly informative given the 54/46 class split. **F1 Score** balances precision and recall; both error directions carry real planning costs, since over-preparing wastes resources and under-preparing degrades participant experience. **5-Fold Cross-Validation Accuracy** is the primary generalisation metric, reported as the mean across five stratified folds.

V. RESULTS AND ANALYSIS

A. Model Performance

Table IV presents the full v1 versus v2 comparison. The most important number is not the 5-point accuracy gain but the 25-point improvement in cross-validation accuracy. In v1, the 20-point gap between test accuracy (83%) and CV accuracy (63%) signalled that the model had not genuinely generalised — it had simply been evaluated on a partition that happened to suit it. In v2, test accuracy and CV accuracy are essentially equal, which means the model is learning stable patterns from the data rather than fitting a convenient split.

TABLE IV
MODEL PERFORMANCE: V1 VS. V2

System	Acc.	AUC-ROC	F1	5-Fold CV
v1 (RF, N=120, 11 feat.)	83%	0.880	0.850	63%
v2 (GB, N=3000, 18 feat.)	**88%**	**0.956**	**0.904**	**88%**
Δ	+5 pp	+0.076	+0.054	+25 pp

An AUC-ROC of 0.956 means the classifier correctly ranks a randomly drawn attendee above a randomly drawn non-attendee in 95.6% of cases — well above the v1 baseline of 0.880 and substantially better than chance. The confusion matrix on the held-out test set shows errors distributed roughly evenly between false positives and false negatives, which matters in practice: a model that is systematically wrong in one direction is far less useful for planning than one whose mistakes are balanced.

B. Comparative Model Analysis

Table V compares all six classifiers evaluated on the v2 feature set. Gradient Boosting leads on both accuracy and AUC-ROC while remaining as compact and fast to deploy as Random Forest. The Neural Network (MLP) approaches Gradient Boosting on accuracy but offers no native feature importance and requires considerably more hyperparameter tuning to reach that level. Logistic Regression is the easiest model to explain to a non-technical audience, but at 76% accuracy it leaves 12 percentage points on the table — a real cost when the output is being used to make resource allocation decisions.

TABLE V
COMPARATIVE MODEL PERFORMANCE ON V2 FEATURE SET (**BOLD** = SELECTED)

Model	Acc.	AUC	Complexity	Interpretability
Linear Regression	72%	0.740	Low	High
Logistic Regression	76%	0.800	Low	High
Decision Tree	79%	0.790	Low	Medium
Random Forest	85%	0.910	Medium	Medium
Gradient Boosting	**88%**	**0.956**	Medium	Medium
Neural Network (MLP)	86%	0.930	High	Low

C. Feature Importance Analysis

Table VI reports the top-ten feature importances computed via Eq. (4).

Three findings stand out from this ranking.

First, the two most important features are both interaction terms, not base features. This confirms that the compound effects the v2 feature layer was designed to capture are

TABLE VI
TOP-10 FEATURE IMPORTANCES (NORMALISED GINI)

Rank	Feature	Imp.	Interpretation
1	`rating_x_org`	0.142	Quality–reputation synergy
2	`buzz_x_rating`	0.118	Engagement amplified by quality
3	`event_rating`	0.097	Standalone quality signal
4	`distance_km`	0.089	Proximity as primary logistics lever
5	`dist_income`	0.076	Income-relative travel barrier
6	`loyalty_score`	0.071	Composite repeat-attendee signal
7	`organizer_score`	0.063	Reputation effect, independent of rating
8	`social_buzz`	0.058	Direct engagement-to-attendance effect
9	`income`	0.051	Income-based attendance conversion
10	`previous_events`	0.047	Repeat-attendance behaviour

real: a well-rated event run by a well-regarded organiser is qualitatively more likely to produce actual attendance than either factor alone suggests. The v1 model, working only with independent inputs, had no way to represent this.

Second, `social_buzz` holds independent importance at rank 8 even after `buzz_x_rating` has already captured its interaction with event quality. This is consistent with Wang et al.'s finding [7] that social visibility has a direct effect on attendance follow-through that operates independently of event quality.

Third, `loyalty_score` (Eq. 1) reaches rank 6 despite being composed of three individually moderate-importance base features. The joint product surfaces a specific subgroup of registrants — those who have attended before, at well-rated events, run by trusted organisers — who behave noticeably differently from the broader pool. Treating these factors as a composite rather than three independent inputs is what allows the model to identify this subgroup cleanly.

D. Ablation Study

To quantify the independent contribution of each feature group, the Gradient Boosting model was retrained on the same 3,000 records with successive groups removed. Table VII reports AUC-ROC under each condition.

The interaction feature layer contributes the largest independent gain: removing it costs 0.030 AUC-ROC, more than either the domain features (0.018) or the temporal features (0.012). This supports the claim that the interaction framework is not a marginal addition but a meaningful transformation of the feature space. It is worth noting that the base-only model (AUC-ROC 0.891) already outperforms v1 (0.880), attributable to the larger dataset — the dataset expansion and the feature engineering are both contributing, and neither alone would have been enough.

TABLE VII
ABLATION STUDY: AUC-ROC BY FEATURE GROUP REMOVAL

Configuration	AUC-ROC	Δ vs. Full
Full model (18 features)	0.956	—
Remove interaction features (7)	0.926	−0.030
Remove domain features (2)	0.938	−0.018
Remove temporal features (3)	0.944	−0.012
Base features only (8)	0.891	−0.065

VI. SYSTEM DEPLOYMENT

EventIQ v2 is deployed as a three-tier system: a SQLite data layer that logs every prediction and builds an institutional audit trail over time, a Flask REST API backend exposing six endpoints (`/health`, `/metadata`, `/predict`, `/analyze`, `/history`, `/stats`), and a Streamlit frontend. A shared `utils/preprocess.py` module handles all feature construction identically at both training and inference time, ensuring consistency across the full pipeline.

The dashboard provides five pages. The **Predict** page is where most organiser interaction happens: a user fills in what they know about a registrant and the system returns a colour-coded probability gauge (red below 40%, amber between 40–60%, green above 60%) alongside a driver bar chart that shows which features pushed the probability up or down for that specific input. This chart makes the model's reasoning legible to anyone reading it, without requiring any understanding of how Gradient Boosting works. The other pages cover aggregate statistics by event type and city (**Dashboard**), model-level feature importances and the confusion matrix (**Insights**), batch prediction from an uploaded registration list (**Analyze CSV**), and interactive exploration of the training data distributions (**EDA**).

VII. DISCUSSION

A. Why the Results Improved

The ablation study makes clear that the improvement from v1 to v2 did not come from a single change. The dataset expansion, the interaction feature layer, and the model upgrade each contributed independently, and removing any one of the three measurably degrades performance. This matters because it means the result is robust: there is no single lucky design decision that the performance depends on.

The dominance of interaction terms in the importance ranking has a practical implication that is easy to state: event quality and organiser reputation do not just add together, they multiply. An organiser who has built a track record of well-run events creates a context that is qualitatively more likely to convert registrations into actual attendance, and the model correctly learns to weight that distinction. For organisers, the takeaway is that investing in programme quality compounds over time — both the rating and the reputation improve together, and both feed the highest-importance features in the model.

B. Limitations

Three limitations are worth stating directly. First, while the use of synthetic data augmentation is standard and well-justified for small institutional datasets, the results on a fully real dataset of equivalent size might differ. Growing the real-data component over time is a practical priority.

Second, `social_buzz` is entered manually at prediction time. Until this is connected to live social media APIs, it relies on the organiser's own assessment of their event's online engagement, which may be optimistic.

Third, the model was trained on data from six urban Indian cities. Applying it without retraining in a substantially different demographic or geographic context — a rural campus, an international conference, a professional training event — would require domain-specific validation first.

VIII. Conclusion and Future Scope

A. Conclusion

This paper presented EventIQ v2, a machine learning system for predicting academic event attendance that is designed to work within the real constraints of institutional settings: small datasets, non-technical users, and limited infrastructure. The central contribution is a seven-term interaction feature framework, motivated by domain-specific behavioural hypotheses, that captures how quality, reputation, social engagement, and logistics combine to determine whether a registrant shows up. An ablation study confirms that this framework contributes an independent AUC-ROC gain of 0.030 beyond what base features alone can achieve. A Gradient Boosting Classifier trained on 3,000 records achieves 88% accuracy and AUC-ROC of 0.956, with five-fold cross-validation confirming stable generalisation.

What this work ultimately demonstrates is that the gap between the data institutions already collect and the planning decisions they make does not require large infrastructure to close. It requires a well-designed pipeline, a feature set grounded in domain knowledge, and a model whose outputs people can understand and act on.

The codebase is publicly available at https://github.com/Ayushaman1291/eventiq-v2.

B. Future Scope

User authentication and role-based access. The current system has no login layer, which limits its use in shared institutional settings. The next version will introduce JWT-based authentication with three roles: administrators who can manage datasets and retrain the model, organisers who can access predictions and history scoped to their own events, and read-only viewers such as faculty coordinators and department heads. This also enables multi-team deployments where different organising groups maintain independent prediction histories.

Automated email delivery. At present, predictions are only accessible within an active dashboard session. We are building an SMTP-based module that will automatically send pre-event summary reports to registered organisers at a configurable number of days before each event, including the predicted attendance rate, the top factors driving the forecast, and suggested logistics adjustments. Batch predictions will also support optional email delivery of results, removing the need for the organiser to be actively logged in.

SHAP-based explanations. The current driver bar chart renders Gini importances relative to each input but does not produce formally grounded attribution values. Replacing it with SHAP (SHapley Additive exPlanations) will make the explanations additive and directly quantified, enabling statements like "the low organiser score reduced this registrant's probability by 12 percentage points" rather than a directional bar without a unit.

Live social media API integration. Automating `social_buzz` by pulling real-time engagement data from social media platforms will remove the manual estimation step and the self-reporting bias it introduces.

Multi-class attendance prediction. Extending the target from binary to an ordinal scale — for example, four bands covering $<25\%$, 25–50%, 50–75%, and $>75\%$ of registered capacity — will enable more granular resource planning decisions.

Automated retraining and containerised deployment. A scheduled pipeline that appends confirmed post-event outcomes and triggers periodic model retraining will allow the system to improve continuously without manual intervention. Docker packaging will make one-command deployment possible across campuses and reduce the technical overhead for institutions deploying the system independently.

References

[1] G. Biau and E. Scornet, "A Random Forest Guided Tour," *TEST*, vol. 25, no. 2, pp. 197–227, 2016.

[2] K. P. Murphy, *Probabilistic Machine Learning: An Introduction*. Cambridge, MA: MIT Press, 2022.

[3] A. Géron, *Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow*, 3rd ed. Sebastopol, CA: O'Reilly Media, 2022.

[4] M. Kuhn and J. Silge, *Tidy Modeling with R*. Sebastopol, CA: O'Reilly Media, 2022. [Online]. Available: https://www.tmwr.org

[5] L. Buitinck *et al.*, "API Design for Machine Learning Software: Experiences from the Scikit-learn Project," *arXiv preprint arXiv:1309.0238v2*, 2023.

[6] G. Shmueli, P. C. Bruce, P. Gedeck, and N. R. Patel, *Data Mining for Business Analytics: Concepts, Techniques, and Applications in Python*. Hoboken, NJ: Wiley, 2020.

[7] S. Qian, T. Zhang, R. Xu, and J. Shao, "Multi-modal Knowledge-aware Event Memory Network for Social Media Posts Popularity Prediction," *IEEE Transactions on Multimedia*, vol. 24, pp. 2206–2219, 2022.

[8] C. Bentéjac, A. Csörgő, and G. Martínez-Muñoz, "A Comparative Analysis of Gradient Boosting Algorithms," *Artificial Intelligence Review*, vol. 54, no. 3, pp. 1937–1967, 2021.

[9] A. Natekin and A. Knoll, "Gradient Boosting Machines, a Tutorial," *Frontiers in Neurorobotics*, vol. 7, p. 21, 2021.

[10] Pallets Projects, "Flask Documentation," 2024. [Online]. Available: https://flask.palletsprojects.com

[11] Streamlit Inc., "Streamlit Documentation," 2024. [Online]. Available: https://docs.streamlit.io